

| Project Name | ... |
|----------------------------|---|
| Online team meeting | https://fau.zoom-x.de/j/3913240515?pwd=TWo2NKZmaVZFWDVKckNKVzQyRG4zdz09 |
| Production system (if any) | TBT |
| Test system (if any) | TBT |
| GitHub repository | https://github.com/amosproj/amos2025ws01-opensearch-load-tester |
| GitHub feature board | https://github.com/orgs/amosproj/projects/89/views/2 |
| GitHub imp-squared backlog | https://github.com/orgs/amosproj/projects/93/views/1 |
| Team T-shirt (white) | https://www.shirtinator.de/en/s/nasAq-u7RjO3FQ4BPPUSLw |
| Team T-shirt (black) | https://www.shirtinator.de/en/s/0VuwTYlwRviJfsV7ERwamA |
| Additional materials | ... |
| Team mailing list | oss-amos-proj1@lists.fau.de |
| | |
| | |

| Last Name | First Name | GitHub User Name | Email Address |
|------------|------------|------------------|--|
| Lorenz | Alexander | Hydraneut | alexander.lorenz@fau.de |
| Buchner | Lea | LeaBuchner | lea.buchner@fau.de |
| Engelhard | Dirk | engelharddirk | dirk.engelhard@fau.de |
| Strachwitz | Carlo | Carlito0 | carlostrachwitz@gmail.com |
| Belz | Sara | SaraBelz | sara.belz@fau.de |
| Becker | Eugen | BeEugen | eugen.becker@fau.de |
| Hofmann | Leo | Leolingio | leohof23@gmail.com |
| Knecht | Sebastian | SeboKnt | sebastian.knecht@fau.de |

| # | Meeting Day | Product Owner | | Software Developer | Release Manager | Scrum Master | Comment |
|---|-------------|----------------|----------------|--------------------|------------------|------------------|-----------------------------|
| | | Review | Planning | | | | |
| 1 | 2025-10-15 | both | both | Everyone else | N/A | Alexander Lorenz | |
| 2 | 2025-10-22 | Lea Buchner | Dirk Engelhard | Everyone else | N/A | Alexander Lorenz | |
| 3 | 2025-10-29 | Dirk Engelhard | Lea Buchner | Everyone else | Carlo Strachwitz | Alexander Lorenz | |
| 4 | 2025-11-05 | Lea Buchner | Dirk Engelhard | Everyone else | Carlo Strachwitz | Alexander Lorenz | Build process review |
| 5 | 2025-11-12 | Dirk Engelhard | Lea Buchner | Everyone else | Leo Hofmann | Alexander Lorenz | |
| 6 | 2025-11-19 | Lea Buchner | Dirk Engelhard | Everyone else | Sara Belz | Alexander Lorenz | |
| 7 | 2025-11-26 | Dirk Engelhard | Lea Buchner | Everyone else | Sebastian Knecht | Alexander Lorenz | Mid-term due |
| 8 | 2025-12-03 | Lea Buchner | Dirk Engelhard | Everyone else | Eugen Becker | Alexander Lorenz | |
| 9 | 2025-12-10 | Dirk Engelhard | Lea Buchner | Everyone else | Eugen Becker | Alexander Lorenz | |
| 10 | 2025-12-17 | Lea Buchner | Dirk Engelhard | Everyone else | Leo Hofmann | Alexander Lorenz | Team Workshop |
| 11 | 2026-01-07 | Dirk Engelhard | Lea Buchner | Everyone else | Sara Belz | Alexander Lorenz | No class but team meeting |
| 12 | 2026-01-14 | Lea Buchner | Dirk Engelhard | Everyone else | Sebastian Knecht | Alexander Lorenz | |
| 13 | 2026-01-21 | Dirk Engelhard | Lea Buchner | Everyone else | Leo Hofmann | Alexander Lorenz | |
| 14 | 2026-01-28 | Lea Buchner | Dirk Engelhard | Everyone else | Sara Belz | Alexander Lorenz | |
| 15 | 2026-02-04 | Dirk Engelhard | Lea Buchner | Everyone else | Sebastian Knecht | Alexander Lorenz | Demo day!+retrospective |
| | | | | | | | |
| Product owners, software developers, and Scurm Master are set and ideally don't change over time; the critical part is the Release Manager role you need to define here | | | | | | | |
| | | | | | | | |

| | |
|--------------------------------|--|
| Goals | achieve the project goal get inside into the scrum and agile process improve practical development skills create something worth using |
| Meeting norms | everybody should feel safe be punctual Absence from meetings should be (if possible) declared a day in advance |
| Working norms | open door policy (everybody should be allowed to do mistakes) maintaining good documentation Backlog items are assigned by the developers themself before they work (outside the team meeting) |
| Coordination norms | every Backlog item has at least 1 dedicated person (if work has already started on it) Contributors assign themselves to backlog items |
| Communication norms | meetings take place in English unless stated otherwise stick to the topic of the channel We use discord as our primary method of communication we check discord at least once a day (and react if needed) We use WhatsApp for urgent communication (less than an hour) |
| Consideration norms | If you see something say something We decide with majority vote If we think we could help we extend it |
| Cont. improvement norms | constructive feedback is always welcome We jointly review the happiness index |
| Rewards | celebrate successes Appreciation and praise |
| Sanctions | apologize sincerely and find a funny way to apologize |
| Signatures | |
| Scrum Master | Alexander Lorenz |
| Product owner | Dirk Engelhard |
| Product owner | Lea Buchner |
| Software developer | Leo Hofmann |
| Software developer | Eugen Becker |
| Software developer | Sebastian Knecht |
| Software developer | Sara Belz |
| Software developer | Carlo Strachwitz |

| Product Vision | Project Mission |
|--|---|
| Make OpenSearch deployments predictable, right-sized, and resilient by providing a standard, evidence-based benchmarking and load-testing platform that reveals real-world limits and behaviors. | Build a modular, extensible load-tester that reliably measures OpenSearch performance limits and produces reproducible artifacts for capacity planning and reliability engineering. |

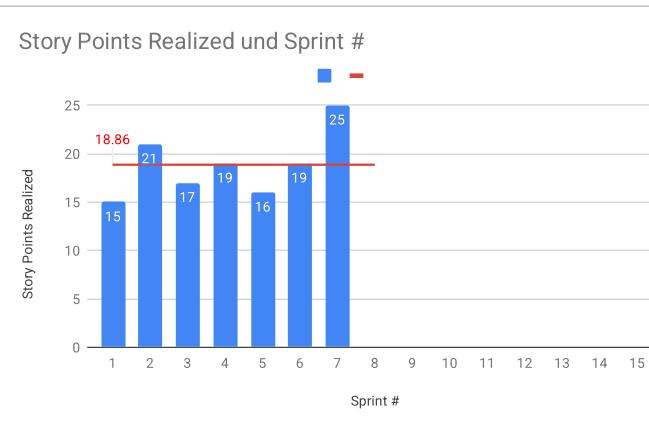
| Term | Definition |
|--------------------------------------|---|
| OpenSearch | A search and analytics engine (Elasticsearch fork). In this project, OpenSearch is the target system for load and ingest testing. |
| Load Generator / load-generator | Component that produces requests/load against OpenSearch to perform performance and stress tests. Located in the load-generator module. |
| Load Test | A planned execution that runs a certain number of queries, a load pattern, or a duration against OpenSearch to measure behavior, throughput and latency. |
| Load Runner / LoadRunnerService | Service component that orchestrates query execution for a test (scheduling, threads, rate limiting, metric collection). |
| Query (Search Query) | A single search or aggregation request, represented as a JSON template under src/main/resources/queries/ (e.g. q1_ano_payroll_range.json). Templates may include parameters/placeholders. |
| Query Template | JSON template for a query with placeholders to be filled at runtime. |
| QueryRunRequest | DTO describing a request to run a query in the load generator (which query, repetitions, rate, etc.). Found in controller/QueryRunRequest.java. |
| Query Execution / QueryExecution | Abstraction for executing a query against a target. Implementations perform the actual call (or a no-op for dry runs). |
| OpenSearchQueryExecution | Concrete QueryExecution implementation that sends requests to OpenSearch. |
| NoOpQueryExecution | A test or dry-run implementation of QueryExecution that does not send real requests. |
| QueryExecutionFactory | Factory producing appropriate QueryExecution instances depending on configuration or target. |
| QueryRegistry | Registry or catalog of available query templates and metadata (names, paths, parameters). |
| Index | OpenSearch index: a logical collection of documents defined by mappings. |
| Document | A single JSON record stored in an index. Test data is created and indexed as documents. |
| Mapping | Schema-definition for an index (fields and types). |
| Aggregation | OpenSearch operation to group or summarize data (counts, sums, bucket aggregations). Examples: q5_ano_clients_aggregation.json. |
| Range Query | A query that filters documents based on numeric or date ranges (e.g., payroll ranges). |
| Match / Term Query | Full-text or exact-match queries in OpenSearch. |
| Bulk API / Bulk Indexing | Mechanism for inserting many documents in a single request to increase throughput. |
| Ingestion Rate | Number of documents indexed per time unit. |
| Throughput | Number of successfully processed requests per second (often QPS). |
| QPS (Queries Per Second) | Metric indicating how many queries are executed per second. |
| Concurrency | Number of parallel requests/threads during a test. |
| Latency | Time between request and response. Typical metrics: average, p95, p99. |
| Response Time | General term for latency; often reported with distribution statistics. |
| Error Rate | Fraction of failed responses (HTTP errors, timeouts) out of total requests. |
| Metrics Collector / metrics-reporter | Component that gathers runtime metrics (latency, throughput, errors) and exports or stores them. See metrics-reporter module. |
| MetricsCollectorService | Service in load-generator that aggregates metrics from test runs and forwards them to a reporter or sink. |
| Test Data | Data that is indexed into OpenSearch for testing, often generated synthetically (see testdata-generator). |
| testdata-generator | Module responsible for generating and optionally storing test datasets (models: AnoRecord, DuoRecord). |

| Term | Definition |
|---|---|
| Ano / AnoRecord | Domain model class in the testdata-generator module. Represents a specific record schema used by some queries (fields defined in the model file). |
| Duo / DuoRecord | Another domain model in testdata-generator with a different schema; used by queries referencing duo data. |
| Recordable | Interface/abstraction implemented by AnoRecord and DuoRecord to provide serialization or indexing behavior. |
| PersistentDataGeneratorService | Service that persists generated test data (e.g., to files or external storage). |
| DynamicDataGeneratorService | Service that generates test data dynamically at runtime (randomized or parameterized fields). |
| FileStorageService | Utility/service for reading and writing generated test data to/from files. |
| OpenSearchDataService | Component that communicates with OpenSearch to index, delete, or query data. |
| Client (business entity) | Business object representing a customer/client. Queries like q7_duo_client_by_customer_number.json reference client data. |
| Booking | Business domain entity (e.g., a financial booking or reservation). Appears in queries like q4_duo_booking_by_client_and_state.json. |
| SLA (Service Level Agreement) - general concept | Target or acceptance values for latency/throughput/error rates used to validate system behavior under load. |
| Dry-run | Mode where queries are not actually sent to OpenSearch (see NoOpQueryExecution). |
| Warmup | Pre-measurement phase in load tests to warm caches and stabilize the system before capturing metrics. |
| Backpressure | System behavior when the target cannot keep up with the requested throughput; relevant to error rates and latency spikes. |

| Sprint # | Sprint goal |
|----------|--|
| 1 | None |
| 2 | None |
| 3 | None |
| 4 | Optional |
| 5 | Internal component integration of each service |
| 6 | First End-To-End test run |
| 7 | Add Warm-Up and Ramp-Up Phase |
| 8 | Query capabilities |
| 9 | Query randomization |
| 10 | Distributed Load Execution |
| 11 | Query Scenarios |
| 12 | Visualization for demo |
| 13 | Stabilization, tuning, packaging |
| 14 | DEMO-DAY readiness |
| 15 | |
| | |
| | |
| | |

| Sprint # | Story Points Realized | Avg |
|---|-----------------------|----------|
| 1 | | 15 18.86 |
| 2 | | 21 18.86 |
| 3 | | 17 18.86 |
| 4 | | 19 18.86 |
| 5 | | 16 18.86 |
| 6 | | 19 18.86 |
| 7 | | 25 18.86 |
| 8 | | 18.86 |
| 9 | | 18.86 |
| 10 | | 18.86 |
| 11 | | 18.86 |
| 12 | | 18.86 |
| 13 | | 18.86 |
| 14 | | 18.86 |
| 15 | | 18.86 |
| PLEASE CREATE THE VELOCITY CHART ON A NEW TAB USING THE DATA FROM THIS TAB | | |
| | | |

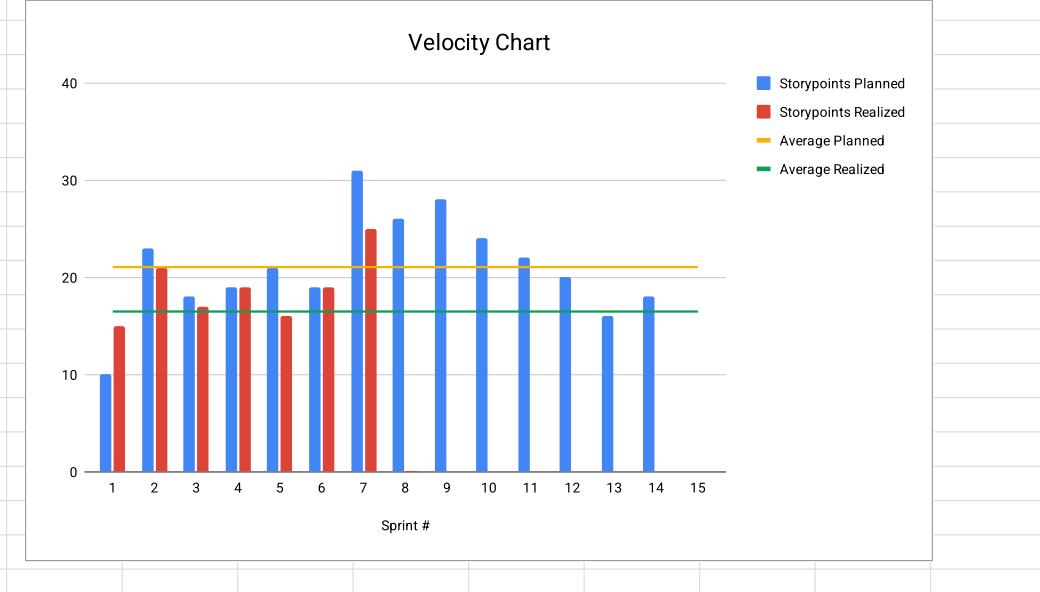
Velocity Chart



Average Speed: 18.86 Points/Sprint

Scrum Master Tables

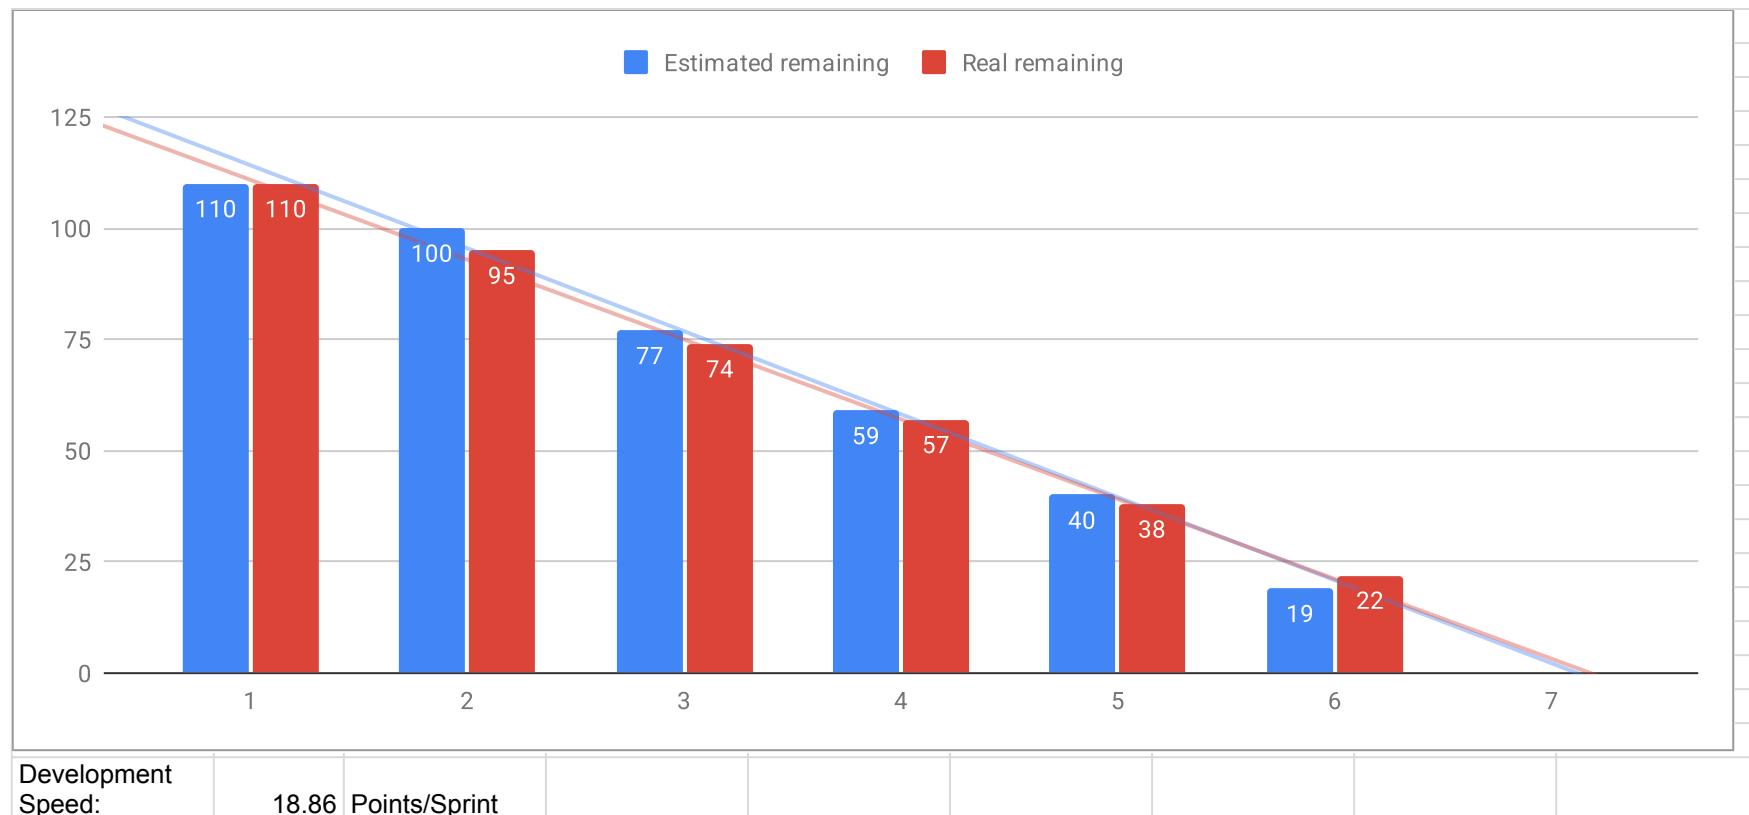
| Sprint | Storypoints Planned | Storypoints Realized | Average Planned | Average Realized |
|--------|---------------------|----------------------|-----------------|------------------|
| 1 | 10 | 15 | 21.07142857 | 16.5 |
| 2 | 23 | 21 | 21.07142857 | 16.5 |
| 3 | 18 | 17 | 21.07142857 | 16.5 |
| 4 | 19 | 19 | 21.07142857 | 16.5 |
| 5 | 21 | 16 | 21.07142857 | 16.5 |
| 6 | 19 | 19 | 21.07142857 | 16.5 |
| 7 | 31 | 25 | 21.07142857 | 16.5 |
| 8 | 26 | 0 | 21.07142857 | 16.5 |
| 9 | 28 | - | 21.07142857 | 16.5 |
| 10 | 24 | - | 21.07142857 | 16.5 |
| 11 | 22 | - | 21.07142857 | 16.5 |
| 12 | 20 | - | 21.07142857 | 16.5 |
| 13 | 16 | - | 21.07142857 | 16.5 |
| 14 | 18 | - | 21.07142857 | 16.5 |
| 15 | - | - | 21.07142857 | 16.5 |



| Sprint | Goal | Feature Name | Est. size | Est. remaining | Real size | Real remaining |
|-----------------|--|--------------|-----------|----------------|-----------|----------------|
| | | | | | | |
| Release | | | | | | |
| Total | | | 110 | 110 | | |
| Sprints | | | | | | |
| 1 | Deliver first increment of initialized software modules and an architecture plan | | 10 | 110 | 15 | 110 |
| 2 | Deliver increment with Build process implementation and first module prototypes | | 23 | 100 | 21 | 95 |
| 3 | Deliver increment with broad prototype implementation coverage | | 18 | 77 | 17 | 74 |
| 4 | Deliver increment with first component integrations implemented | | 19 | 59 | 19 | 57 |
| 5 | Deliver increment with more integration and technical test running capability | | 21 | 40 | 16 | 38 |
| 6 | Deliver Increment with simple full test run routine executable | | 19 | 19 | 19 | 22 |
| 7 | | | | | | |
| Features | | | | | | |
| 1 | Deliver first increment of initialized software modules and an architecture plan | | | | | |
| | Research - OpenSearch | | 3 | | 3 | |
| | Research - Understand DATEV Data Structure (Industry Partner Input) for Load Tester Query Design | | 1 | | 1 | |
| | Developer Environment Setup | | 1 | | 1 | |
| | GitHub Repository Management Decision | | | | 2 | |
| | Architecture evaluation/decision | | 2 | | 5 | |
| | Design Team Shirt | | 1 | | 1 | |
| | Initialize SBOM | | | | | |
| | Minimal Module Initialization | | 2 | | 2 | |
| 2 | Deliver increment with Build process implementation and first module prototypes | | | | | |
| | Define Representative Query Scenarios | | 5 | | 3 | |
| | Prepare Build Process Review | | 2 | | 2 | |
| | Research - Benchmark Tools for OpenSearch Load Tester | | 3 | | 3 | |
| | OpenSearchClient prototype | | 3 | | 3 | |
| | Build Process Decision | | 2 | | 2 | |
| | Data Generation Prototype | | 3 | | 3 | |
| | Architecture adaptation | | 2 | | 2 | |
| | Build Process Implementation | | 2 | | 2 | |
| | Github Signoff/Co-Author Automation | | 1 | | 1 | |
| 3 | Deliver increment with broad prototype implementation coverage | | | | | |
| | MetricsCollector Prototype Implementation | | 5 | | 3 | |

| Sprint | Goal | Feature Name | Est. size | Est. remaining | Real size | Real remaining |
|----------|---|---|-----------|----------------|-----------|----------------|
| | | QueryExecution Prototypee Implementation | 3 | | 2 | |
| | | LoadRunner Component Prototype Implementation | 2 | | 2 | |
| | | Logging Mechanism Implementation | 2 | | 2 | |
| | | Data Generation <> OpenSearch Client Integration | 3 | | 5 | |
| | | Environment Variable Support for Dockerized Services | 2 | | 2 | |
| | | Automatic Shutdown of Testdata Generator Service | 1 | | 1 | |
| | | | | | | |
| 4 | Deliver increment with first component integrations implemented | | | | | |
| | | MetricsCollector Prototype Implementation | 5 | | 3 | |
| | | QueryExecution Prototype Implementation | 3 | | 2 | |
| | | LoadRunner Component Prototype Implementation | 2 | | 2 | |
| | | Data Generation with OpenSearch Client Integration | 3 | | 5 | |
| | | Logging Mechanism Implementation | 2 | | 2 | |
| | | Implementation of Batch Processing for Data Pre-Loading | 3 | | | |
| | | Unified Code Styling & Editor Configurations | 1 | | 2 | |
| | | Ano/Duo Separation | 3 | | 3 | |
| | | | | | | |
| 5 | Deliver Increment with internal component integrations broadly implemented | | | | | |
| | | Load Generator Port Setup Dockerized Services | 3 | | 2 | |
| | | Load Generator Internal Component Integration | 5 | | 5 | |
| | | ReportController Component Prototype Implementation | 5 | | 3 | |
| | | Build Process Video | 2 | | 2 | |
| | | Definition of Done | 3 | | 2 | |
| | | Implementation of MetricsReporterClient Prototype | 3 | | 2 | |
| | | | | | | |
| 6 | Deliver Increment with simple full test run routine executable | | | | | |
| | | Implementation of Load Test Scenario | 5 | | 5 | |
| | | Implementation of Batch Processing for Data Pre-Loading | 3 | | 3 | |
| | | ReportCreator Component Prototype Implementation | 5 | | 8 | |
| | | Wiki Creation | 3 | | 3 | |
| | | CI/CD Pipeline Setup | 3 | | | |

Mid-Project Burndown Chart

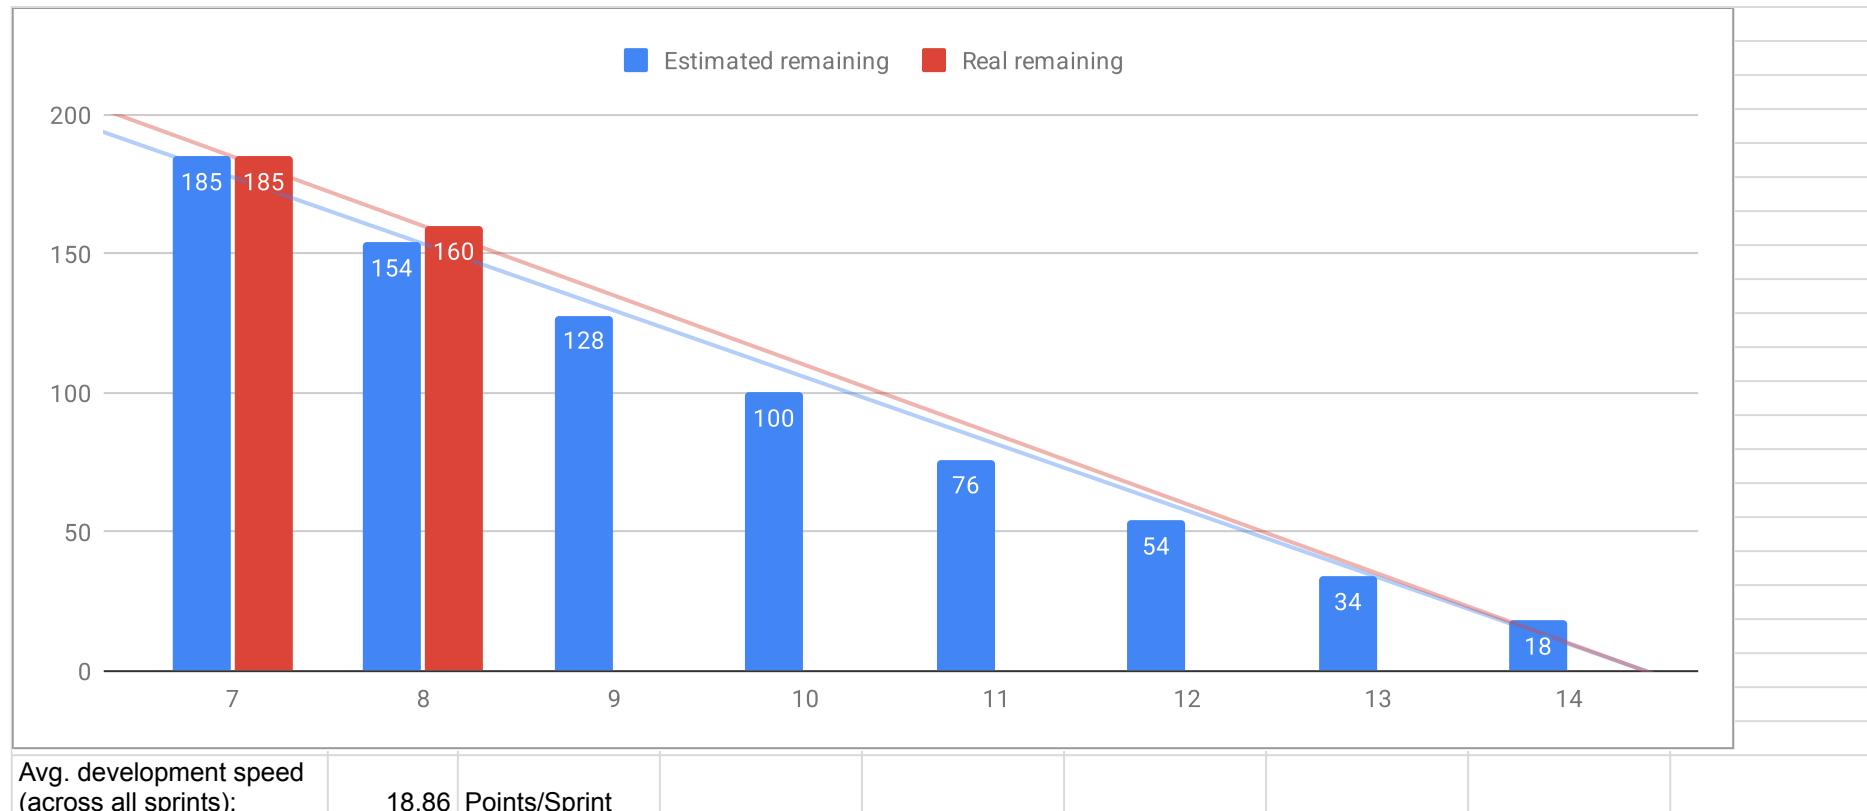


| Sprint | Goal | Feature Name | Est. size | Est. remaining | Real size | Real remaining |
|-----------------|---|--------------|-----------|----------------|-----------|----------------|
| | | | | | | |
| Release | | | | | | |
| Total | | | 185 | 185 | | |
| Sprints | | | | | | |
| 7 | Deliver Increment with full test run executable supporting warm-up and synchronized load start | | 31 | 185 | 25 | 185 |
| 8 | Deliver Increment with improved query capabilities | | 26 | 154 | 0 | 160 |
| 9 | Deliver Increment with broad query randomization | | 28 | 128 | | |
| 10 | Deliver Increment with Distributed load execution | | 24 | 100 | | |
| 11 | Deliver Increment with set of query scenarios | | 22 | 76 | | |
| 12 | Deliver Increment with a data report visualization for demo purposes | | 20 | 54 | | |
| 13 | Deliver Increment with Stabilization, tuning, packaging | | 16 | 34 | | |
| 14 | Deliver Increment with finished, DEMO-DAY-Ready product | | 18 | 18 | | |
| | | | | 100 | | 0 |
| | | | | | | |
| Features | | | | | | |
| 7 | Deliver Increment with full test run executable supporting warm-up and synchronized load start | | | | | |
| | Complex Queries | | 3 | | | |
| | Big Data Handling in MetricsReporter | | 5 | | 5 | |
| | Code Refactoring | | 5 | | 5 | |
| | Shared Library Creation | | 1 | | 1 | |
| | Configurable Queries per Second | | 3 | | 5 | |
| | Load Generator Request Synchronization | | 5 | | | |
| | Optimize Random Generation of Test Data | | 3 | | 3 | |
| | Configurable Test Duration | | 3 | | 3 | |
| | Load Generator functionality Definition | | 3 | | 3 | |
| | | | | | | |
| 8 | Deliver Increment with improved query capabilities | | | | | |
| | Redesign Metrics Data Model | | 2 | | | |
| | Extend Ano and Duo Index Mapping and Settings | | 3 | | | |
| | Refactoring: Include DTO changes in Metrics Reporter | | 2 | | | |
| | Query Parameter Randomization | | 5 | | | |
| | Warm-Up Phase | | 1 | | | |
| | Implementation of Batch Processing for Data Pre-Loading | | 3 | | | |
| | Documentation for Advanced Query Usage | | 2 | | | |

| Sprint | Goal | Feature Name | Est. size | Est. remaining | Real size | Real remaining |
|--------|--|---|-----------|----------------|-----------|----------------|
| | | Load Generator Start Synchronization | 3 | | | |
| | | Extend (Unit) Test Coverage | 5 | | | |
| | | Test ending wait mechanism | 2 | | | |
| | | CI/CD Pipeline Setup | 2 | | | |
| | | Complex Queries | 2 | | | |
| 9 | Deliver Increment with broad query randomization | | | | | |
| | | Load Generator Query Type Mix | 3 | | | |
| | | Query Refactoring | 5 | | | |
| | | Architecture (Diagram) Refactoring | 3 | | | |
| | | Compound Queries | 5 | | | |
| | | Leaf Queries: specialized/span queries | 5 | | | |
| | | Expensive Queries | 5 | | | |
| 10 | Deliver Increment: Distributed load execution | | | | | |
| | | Multi-Node Load Generation Support | 5 | | | |
| | | Coordinator Node Architecture | 3 | | | |
| | | Load Distribution Algorithm | 3 | | | |
| | | Resilient Messaging Between Nodes | 5 | | | |
| | | Distributed Metrics Aggregation | 3 | | | |
| 11 | Deliver Increment: Set of query scenarios | | | | | |
| | | Define Scenario Structure & DSL | 5 | | | |
| | | Pre-built Scenarios (Search, Aggregations, Mixed) | 5 | | | |
| | | Time-based Scenario Control | 3 | | | |
| | | Scenario Looping & Weighted Groups | 3 | | | |
| | | Scenario Summary & Reporting | 3 | | | |
| 12 | Deliver Increment: Data report visualization for demo | | | | | |
| | | Basic Web Dashboard | 3 | | | |
| | | Real-time Graphs (QPS, Latency, Errors) | 5 | | | |
| | | Trend Lines & Historical Data Comparison | 3 | | | |
| | | Export Reports (PDF/CSV) | 3 | | | |
| | | Dashboard Mobile-Friendly UI | 2 | | | |
| | | Demo Visualization Themes | 2 | | | |
| 13 | Deliver Increment: Stabilization, tuning, packaging | | | | | |
| | | Load Generator Optimization (GC, Memory) | 5 | | | |

| Sprint | Goal | Feature Name | Est. size | Est. remaining | Real size | Real remaining |
|-----------|--|----------------------------------|-----------|----------------|-----------|----------------|
| | | Query Execution Tuning | 3 | | | |
| | | Packaging into Standalone Binary | 3 | | | |
| | | End-to-End Stability Tests | 5 | | | |
| | | Documentation Cleanup | 2 | | | |
| | | | | | | |
| 14 | Deliver Increment with finished, DEMO-DAY-Ready product | | | | | |
| | | Final Demo Workflow | 3 | | | |
| | | Polished Dashboard & UX | 5 | | | |
| | | Scenario Library Complete | 3 | | | |
| | | Full System QA | 5 | | | |
| | | Presentation + Demo Script | 5 | | | |
| | | | | | | |

Final Project Burndown Chart



| # | Feature Definition of Done | Sprint Release Definition of Done | Project Release Definition of Done |
|---------------------------------|---|--|---|
| Overview | A single new feature or user story implemented in Trunk Load (e.g., new query type, new CLI option, new scenario). | The subset of functionality and artifacts that must be finished and deliverable at the end of a sprint (sprint-level deliverable). | The final project release: full handover of the stress-testing toolset, reproducible scenarios, baselines and deliverables. |
| Scope | Code, tests and docs needed to implement the feature in load-generator, metrics-reporter or testdata-generator. | All features/bugfixes scoped to the sprint board's "Done" column, plus required scenarios and deliverables in Deliverables/sprint-XX/. | All features and fixes required for the final deliverable, automated scenarios (ramp, steady, spike, soak, chaos), reporting assets and reproducibility instructions in Deliverables/. |
| Acceptance Criteria | <ul style="list-style-type: none"> - Implements the story's acceptance criteria from the issue. - At least one concrete example scenario demonstrating the feature exists. - No secrets added and no new critical static-analysis findings. - Excludes wide refactors unless explicitly requested. - Code compiles without errors locally and in CI. - Formatting and linting checks pass. - Minimal unit tests cover the relevant logic. - Code is documented where appropriate (Wiki/internal README). | <ul style="list-style-type: none"> - All sprint issues marked Done have merged PRs. - Each merged item satisfies its Feature or Bugfix DoD. - All completed features are integrated and do not break the full system. - Architectural guidelines are respected. - Full Docker Compose setup starts without errors: <ul style="list-style-type: none"> o OpenSearch + Dashboards o Load Generator o Metrics Reporter / Metrics Storage (if present) o Report Controller o Logging stack (Promtail/Loki) if used o Metrics Reporter / Metrics Storage (if present) - The increment can be demonstrated in the Sprint Review without failures. | <ul style="list-style-type: none"> - All required scenarios implemented and automated. - Reproducible runs produce archived results (raw + aggregated charts). - Required performance/resilience targets met for mandatory scenarios. - Documentation is complete (test plan, runbook, interpretation guide, reproducibility checklist). - No open critical bugs. - All promised components implemented: <ul style="list-style-type: none"> o Fully working Load Generator o Query Execution + Templates o Metrics Reporter o Report Controller (prototype level is OK if IP handles the rest) - Stable OpenSearch integration - Data generation for "ano" and "duo" - Final code review approved by the team. - No critical warnings or dead code. - Repository clean with no unmerged branches. |
| Tests & Validation | <ul style="list-style-type: none"> - Unit tests: new/changed classes covered (place under module src/test/java). - Integration test: Testcontainers or short docker-compose integration that exercises the feature against OpenSearch. - Smoke run: a short end-to-end scenario (e.g., 30–60s) that demonstrates feature behavior and produces output. - Basic manual testing completed. - Logs checked and no unexpected errors appear. - Code review completed and approved by at least one teammate. - DTOs, controllers, YAML config files updated as needed | <ul style="list-style-type: none"> (- Run the sprint acceptance scenarios: population → run → collect → validate.) (- Validation scripts assert basic pass/fail (no critical errors, p95/p99 within sprint thresholds).) All commits pass CI: <ul style="list-style-type: none"> o Build o Unit tests o Static analysis (lint) o Formatting - No critical errors in logs. | <ul style="list-style-type: none"> - Full end-to-end validation runs executed (including at least one soak and one chaos test) and validated by scripts. - Regression/reproducibility check: repeat a representative scenario and verify results within accepted variance (e.g., ±10%). - Unit + integration suites green. - Docker Compose environment stable and reproducible. - Entire stack starts from scratch without errors. - Final load tests executed and documented. - Logs structured and free of critical problems. - Final demonstration with DATEV/IP completed. - Test coverage includes: <ul style="list-style-type: none"> o Happy-path scenarios o Error handling (e.g., OS down, timeout, invalid configs, multi-threading) - Minimal stress tests documented. |
| Documentation | - Usage example added to module README.md and/or wiki. | <ul style="list-style-type: none"> - Sprint notes updated in the wiki. - Architecture diagrams updated if anything changed. - API documentation updated with new endpoints. | <ul style="list-style-type: none"> - Deliverables/ contains final reports, scenario definitions, scripts to reproduce, dashboard exports and a README with exact reproduction commands. - Documentation/ contains runbook, interpretation guide, and troubleshooting steps, final architecture, docker setup, API endpoints, execution examples, "How to add new queries", "how to generate and upload data to OpenSearch" as part of wiki. - Final presentation delivered |
| CI / Automation (when existing) | <ul style="list-style-type: none"> - PR CI must run build + unit tests + linter. - Integration smoke test recommended in merge pipeline (if it's fast); otherwise on main/merge job. | <ul style="list-style-type: none"> - Merge/main pipeline runs smoke acceptance scenarios (short) or a job that can be manually triggered to run them. - CI artifacts for the sprint (result JSON/CSV and summary) are archived in pipeline artifacts. | <ul style="list-style-type: none"> - Release pipeline builds artifacts (JARs, Docker images) and archives results. - Nightly or release-run perf jobs executed and artifacts persisted; CI flags regressions compared to baselines. - Dependency/security scans completed. |
| Metrics & Thresholds | <ul style="list-style-type: none"> - If performance-related: attach baseline and verify no significant regression (>10% on critical metric) unless approved. - Smoke validation: error rate < 5% (adjustable per scenario) and p95 reported. <p>PR Checklist</p> <ul style="list-style-type: none"> - Issue/story linked - Unit tests added + passing - Integration/smoke test added - Documentation updated - Reviewer approved | <ul style="list-style-type: none"> - Sprint acceptance thresholds defined per scenario (e.g., for that sprint: p95 <= 600 ms, error rate <= 2%). - If thresholds not met, the sprint cannot be marked finished until remediation or explicit acceptance by PO. <p>PR Checklist</p> <ul style="list-style-type: none"> - All sprint issues linked and merged - Acceptance scenarios included and runnable - Deliverables/sprint-XX populated with results - CI artifacts archived - PO/TA acknowledged | <ul style="list-style-type: none"> - Project-level mandatory thresholds must be met (define concretely before final validation, e.g.): <ul style="list-style-type: none"> - Steady-state p95 <= X ms, p99 <= Y ms, error rate <= Z%. - Soak: sustained stability for N minutes without memory growth or node crashes. - Chaos: cluster remains available and recovers within R minutes. - Reproducibility: repeated runs within ±10% of throughput and latency targets (tunable). <p>PR / Release Checklist</p> <ul style="list-style-type: none"> - All scenarios automated and in repo - Deliverables folder contains final reports and artifacts - Release build artifacts created (JARs, Docker images) - Baselines and comparison scripts in place - Security/dependency checks done - PO/TA and technical reviewers sign off |

Definition of Done

| # | Context | Name | Version | License | Comment |
|----|-------------|---|----------|----------------------------------|---|
| 1 | environment | Eclipse Temurin OpenJDK | 25 (LTS) | GPL-2.0 with Classpath Exception | Java runtime used for building and running the application |
| 2 | environment | Docker | 28.5.1 | Apache-2.0 | Container environment for running OpenSearch and Load Tester services |
| 3 | tool | Apache Maven | 3.9.11 | Apache-2.0 | Build automation and dependency management tool |
| 4 | library | spring-boot-starter | 3.5.7 | Apache-2.0 | Entry point for Spring Boot dependency management |
| 5 | library | spring-boot | 3.5.7 | Apache-2.0 | Core framework providing auto-configuration and runtime support |
| 6 | library | spring-context | 6.2.12 | Apache-2.0 | Dependency injection and application context management |
| 7 | library | spring-aop | 6.2.12 | Apache-2.0 | Aspect-oriented programming support |
| 8 | library | spring-beans | 6.2.12 | Apache-2.0 | Bean creation and configuration framework |
| 9 | library | spring-expression | 6.2.12 | Apache-2.0 | Spring Expression Language (SpEL) processing |
| 10 | library | spring-boot-autoconfigure | 3.5.7 | Apache-2.0 | Automatically configures components based on classpath |
| 11 | library | spring-boot-starter-logging | 3.5.7 | Apache-2.0 | Default logging configuration using Logback |
| 12 | library | logback-classic | 1.5.20 | EPL-1.0 | Logging backend implementing SLF4J API |
| 13 | library | logback-core | 1.5.20 | EPL-1.0 | Core utilities for Logback logging |
| 14 | library | log4j-to-slf4j | 2.24.3 | Apache-2.0 | Redirects Log4j 2 logs to SLF4J |
| 15 | library | jul-to-slf4j | 2.0.17 | MIT | Bridges java.util.logging to SLF4J |
| 16 | library | jakarta.annotation-api | 2.1.1 | EPL-2.0 | Jakarta annotations used by Spring components |
| 17 | library | spring-core | 6.2.12 | Apache-2.0 | Core utilities and classloading framework of Spring |
| 18 | library | spring-jcl | 6.2.12 | Apache-2.0 | Logging abstraction used internally by Spring |
| 19 | library | snakeyaml | 2.4 | Apache-2.0 | YAML parser for configuration files |
| 20 | library | micrometer-core | 1.15.5 | Apache-2.0 | Core metrics collection API |
| 21 | library | micrometer-commons | 1.15.5 | Apache-2.0 | Shared utilities for metrics instrumentation |
| 22 | library | micrometer-observation | 1.15.5 | Apache-2.0 | Captures timing and observation data for metrics |
| 23 | library | HdrHistogram | 2.2.2 | CC0-1.0 | High-resolution latency measurement library |
| 24 | library | LatencyUtils | 2.0.3 | CC0-1.0 | Helper utilities for latency tracking |
| 25 | library | micrometer-registry-prometheus | 1.15.5 | Apache-2.0 | Exposes Micrometer metrics to Prometheus |
| 26 | library | prometheus-metrics-core | 1.3.10 | Apache-2.0 | Core Prometheus metrics collection |
| 27 | library | prometheus-metrics-model | 1.3.10 | Apache-2.0 | Data model for Prometheus metrics |
| 28 | library | prometheus-metrics-config | 1.3.10 | Apache-2.0 | Configuration utilities for Prometheus metrics |
| 29 | library | prometheus-metrics-tracer-common | 1.3.10 | Apache-2.0 | Common tracing utilities for Prometheus exporters |
| 30 | library | prometheus-metrics-exposition-formats | 1.3.10 | Apache-2.0 | Serialization formats for Prometheus metrics |
| 31 | library | prometheus-metrics-exposition-textformats | 1.3.10 | Apache-2.0 | Text exposition format for Prometheus |

| # | Context | Name | Version | License | Comment |
|----|---------|-----------------------------------|---------|--------------|--|
| 32 | library | lombok | 1.18.34 | MIT | Generates boilerplate code (getters/setters, builders, etc.) |
| 33 | library | opensearch-rest-high-level-client | 2.16.0 | Apache-2.0 | REST client for interacting with OpenSearch clusters |
| 34 | library | opensearch | 2.16.0 | Apache-2.0 | Core OpenSearch library |
| 35 | library | opensearch-common | 2.16.0 | Apache-2.0 | Common interfaces and utilities for OpenSearch modules |
| 36 | library | opensearch-core | 2.16.0 | Apache-2.0 | Core APIs for OpenSearch operations |
| 37 | library | opensearch-compress | 2.16.0 | Apache-2.0 | Compression utilities used internally by OpenSearch |
| 38 | library | zstd-jni | 1.5.5-5 | BSD-2-Clause | Java bindings for Zstandard compression |
| 39 | library | opensearch-secure-sm | 2.16.0 | Apache-2.0 | Security manager integration for OpenSearch |
| 40 | library | opensearch-x-content | 2.16.0 | Apache-2.0 | Serialization/deserialization utilities in OpenSearch |
| 41 | library | jackson-dataformat-smile | 2.19.2 | Apache-2.0 | Jackson module for Smile binary JSON |
| 42 | library | jackson-dataformat-cbor | 2.19.2 | Apache-2.0 | Jackson module for CBOR binary JSON |
| 43 | library | opensearch-geo | 2.16.0 | Apache-2.0 | Geospatial data and query support for OpenSearch |
| 44 | library | opensearch-telemetry | 2.16.0 | Apache-2.0 | Telemetry and monitoring features for OpenSearch |
| 45 | library | lucene-core | 9.11.1 | Apache-2.0 | Core search and indexing engine powering OpenSearch |
| 46 | library | lucene-analysis-common | 9.11.1 | Apache-2.0 | Text analysis and tokenization utilities |
| 47 | library | lucene-backward-codecs | 9.11.1 | Apache-2.0 | Backward-compatible codecs for old Lucene versions |
| 48 | library | lucene-grouping | 9.11.1 | Apache-2.0 | Implements result grouping and field collapsing |
| 49 | library | lucene-highlighter | 9.11.1 | Apache-2.0 | Highlighting support for search result snippets |
| 50 | library | lucene-join | 9.11.1 | Apache-2.0 | Supports join queries between documents |
| 51 | library | lucene-memory | 9.11.1 | Apache-2.0 | In-memory index structures for fast search |
| 52 | library | lucene-misc | 9.11.1 | Apache-2.0 | Miscellaneous utilities and experimental features |
| 53 | library | lucene-queries | 9.11.1 | Apache-2.0 | Advanced query utilities and filters |
| 54 | library | lucene-queryparser | 9.11.1 | Apache-2.0 | Parses textual queries into Lucene queries |
| 55 | library | lucene-sandbox | 9.11.1 | Apache-2.0 | Experimental Lucene modules |
| 56 | library | lucene-spatial-extras | 9.11.1 | Apache-2.0 | Spatial/geographic query utilities |
| 57 | library | lucene-spatial3d | 9.11.1 | Apache-2.0 | 3D geospatial shapes and indexing |
| 58 | library | lucene-suggest | 9.11.1 | Apache-2.0 | Suggestion/autocomplete engine |
| 59 | library | opensearch-cli | 2.16.0 | Apache-2.0 | Command-line tools for OpenSearch management |
| 60 | library | jopt-simple | 5.0.4 | MIT | Command-line option parser |
| 61 | library | joda-time | 2.12.7 | Apache-2.0 | Date and time handling library |
| 62 | library | t-digest | 3.2 | Apache-2.0 | Statistical distribution estimation for percentiles |
| 63 | library | log4j-api | 2.24.3 | Apache-2.0 | Core Log4j 2 API |

| # | Context | Name | Version | License | Comment |
|----|---------|--------------------------|---------|------------------------------------|---|
| 64 | library | log4j-jul | 2.24.3 | Apache-2.0 LGPL-2.1-or-later OR | Bridge for Java Util Logging to Log4j |
| 65 | library | jna | 5.13.0 | Apache-2.0 | Java Native Access for system-level operations |
| 66 | library | jzlib | 1.1.3 | BSD-4-Clause | Java implementation of zlib compression |
| 67 | library | reactor-core | 3.7.12 | Apache-2.0 | Reactive programming support (Project Reactor) |
| 68 | library | reactive-streams | 1.0.4 | MIT-0 | Minimal reactive streams specification |
| 69 | library | protobuf-java | 3.22.3 | BSD-3-Clause | Google Protocol Buffers serialization |
| 70 | library | opensearch-rest-client | 2.16.0 | Apache-2.0 | Low-level REST client for OpenSearch |
| 71 | library | httpclient | 4.5.14 | Apache-2.0 | Apache HTTP client for network communication |
| | | | | | Core transport components for Apache HTTP client |
| 72 | library | httpcore | 4.4.16 | Apache-2.0 | |
| 73 | library | httpasyncclient | 4.1.5 | Apache-2.0 | Asynchronous HTTP client for non-blocking I/O |
| 74 | library | httpcore-nio | 4.4.16 | Apache-2.0 | NIO transport layer for asynchronous HTTP |
| 75 | library | commons-codec | 1.18.0 | Apache-2.0 | Utility library for encoding/decoding formats |
| 76 | library | commons-logging | 1.2 | Apache-2.0 | Lightweight logging abstraction |
| 77 | library | mapper-extras-client | 2.16.0 | Apache-2.0 | Additional mapping features for OpenSearch |
| 78 | library | parent-join-client | 2.16.0 | Apache-2.0 | Supports parent-child relationships in OpenSearch |
| 79 | library | aggs-matrix-stats-client | 2.16.0 | Apache-2.0 | Advanced aggregation statistics module |
| 80 | library | rank-eval-client | 2.16.0 | Apache-2.0 | Ranking evaluation tools for OpenSearch |
| 81 | library | lang-mustache-client | 2.16.0 | Apache-2.0 | Template rendering support (Mustache) |
| 82 | library | compiler | 0.9.14 | Apache-2.0 | Annotation processing and compilation utilities |
| 83 | library | jackson-dataformat-yaml | 2.19.2 | Apache-2.0 | YAML serialization and deserialization |
| 84 | library | jackson-databind | 2.19.2 | Apache-2.0 | Object-JSON mapping framework |
| | | | | | Annotations for Jackson serialization/deserialization |
| 85 | library | jackson-annotations | 2.19.2 | Apache-2.0 | |
| 86 | library | jackson-core | 2.19.2 | Apache-2.0 | Core streaming JSON processor |
| 87 | library | slf4j-api | 2.0.17 | MIT | Unified logging abstraction layer |

| Last Name | First Name | Value | | | |
|--|------------|-------|----|------------------|--|
| Lorenz | Alexander | 5 | | | |
| Buchner | Lea | 5 | | | |
| Engelhard | Dirk | | | | |
| Strachwitz | Carlo | | | | |
| Belz | Sara | | 0 | No size | |
| Becker | Eugen | | 1 | Trivial size | |
| Hofmann | Leo | | 2 | Small size | |
| Knecht | Sebastian | | 3 | Medium size | |
| | | | 5 | Large size | |
| | | | 8 | Very large size | |
| | | | 13 | Too large (size) | |
| How to play planning poker | | | | | |
| 1. Everyone type their number into their value field, don't hit return yet | | | | | |
| 2. Someone, perhaps a product owner, count down 3.. 2.. 1.. | | | | | |
| 3. Then, everyone hit return to submit their value | | | | | |